

# Regression interpretation 3 - 4


Oh SukJu

## Contents

0.1	3. Diagnostics for Simple Linear Regression . . . . .	1
0.2	4. Multiple linear regression . . . . .	15

### 0.1 3. Diagnostics for Simple Linear Regression

```
# page 10
set.seed(101)
X = seq(0, 20, length = 21)
Y = 0.5 * X + 1 + rnorm(21)
Y.lm = lm(Y ~ X)
meanY = mean(Y)
Yhat = predict(Y.lm)
plot(X,Y, pch=23, bg='red', cex=2)
```



```
# page 11

plot(X, Y, pch = 23, bg = "red", main='Total sum of squares', cex=2)


abline(Y.lm, pch=23, col='green', lwd=2)

abline(coef(Y.lm), pch=23, col='darkblue', lwd=2)

abline(h = meanY, col = "yellow", lwd = 2)

for (i in 1:21) {
  points(X[i], meanY, pch = 23, bg = "yellow")
  lines(c(X[i], X[i]), c(Y[i], meanY))
}
```

## Total sum of squares



```
# page 12

plot(X, Y, pch = 23, bg = "red", main="Error sum of squares", cex=2)

abline(Y.lm, col = "green", lwd = 2)

for (i in 1:21) {


  points(X[i], Yhat[i], pch = 23, bg = "green")

  lines(c(X[i], X[i]), c(Y[i], Yhat[i]))

}

abline(h = meanY, col = "yellow", lwd = 2)
```

## Error sum of squares



```
# page 13

plot(X, Y, pch = 23, bg = "red", main="Regression sum of squares", cex=2)

abline(Y.lm, col = "green", lwd = 2)

abline(h = meanY, col = "yellow", lwd = 2)

for (i in 1:21) {


  points(X[i], Yhat[i], pch = 23, bg = "green")

  points(X[i], meanY, pch = 23, bg = "yellow")

  lines(c(X[i], X[i]), c(meanY, Yhat[i]))

}
```

## Regression sum of squares



*# page 14*

```
load("/Users/ohsukju/Downloads/data/wages.Rdata")
```

```
wages.lm = lm(logwage ~ education, data=wages)
```

```
SSE = sum(resid(wages.lm)^2)
```

```
SST = sum((wages$logwage - mean(wages$logwage))^2)
```

```
SSR = sum((mean(wages$logwage) - predict(wages.lm))^2)
```

```
data.frame(SST, SSE + SSR)
```

```
##          SST SSE...SSR
```

```
## 1 410.2148 410.2148
```

```
F = (SSR / 1) / (SSE / wages.lm$df)
```

```
print(F)
```

```
## [1] 340.0297
```

```
summary(wages.lm)
```

```
##
## Call:
## lm(formula = logwage ~ education, data = wages)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.78239 -0.25265  0.01636  0.27965  1.61101
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.239194   0.054974   22.54  <2e-16 ***
## education    0.078600   0.004262   18.44  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4038 on 2176 degrees of freedom
## Multiple R-squared:  0.1351, Adjusted R-squared:  0.1347
## F-statistic:   340 on 1 and 2176 DF,  p-value: < 2.2e-16
```

```
2176*0.1351 / (1 - 0.1351)
```

```
## [1] 339.8978
```

```
cor(wages$education, wages$logwage)^2
```

```
## [1] 0.1351453
```

```
# page 17
```

```
summary(wages.lm)
```

```
##
## Call:
## lm(formula = logwage ~ education, data = wages)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.78239 -0.25265  0.01636  0.27965  1.61101
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.239194   0.054974   22.54  <2e-16 ***
## education    0.078600   0.004262   18.44  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4038 on 2176 degrees of freedom
## Multiple R-squared:  0.1351, Adjusted R-squared:  0.1347
## F-statistic:   340 on 1 and 2176 DF,  p-value: < 2.2e-16
```

```
18.44**2
```

```
## [1] 340.0336
```

```
# page 20
```

```
summary(wages.lm)
```

```
##
## Call:
## lm(formula = logwage ~ education, data = wages)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.78239 -0.25265  0.01636  0.27965  1.61101
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.239194   0.054974   22.54  <2e-16 ***
## education    0.078600   0.004262   18.44  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4038 on 2176 degrees of freedom
## Multiple R-squared:  0.1351, Adjusted R-squared:  0.1347
## F-statistic:   340 on 1 and 2176 DF,  p-value: < 2.2e-16
```

```
qf(0.95, 1, 2176)
```

```
## [1] 3.845736
```

```
# page 23
```

```
load("/Users/ohsukju/Downloads/data/anscombe.Rdata")
```

```
y = anscombe$y2
```


```
x = anscombe$x2
```

```
y = y + rnorm(length(y)) * 0.45
```

```
plot(x, y, pch = 23, bg = "orange", cex = 2, ylab = "Y", xlab = "X")
```


```
simple.lm = lm(y ~ x)
```

```
abline(simple.lm, lwd = 2, col = "red", lty = 2)
```



```
# page 24  
plot(x, resid(simple.lm), ylab = "Residual", xlab = "X", pch = 23, bg = "orange", cex = 2)  
abline(h = 0, lwd = 2, col = "red", lty = 2)
```






```
# page 25

x2<-x^2


quadratic.lm = lm(y ~ poly(x, 2))
quadratic.lm2 = lm(y ~ x+x2)
Xsort = sort(x)

plot(x, y, pch = 23, bg = "orange", cex = 2, ylab = "Y", xlab = "X")
lines(Xsort, predict(quadratic.lm, list(x = Xsort)), col = "red", lty = 2, lwd = 2)
lines(Xsort, predict(quadratic.lm2, list(x = Xsort, x2=Xsort^2)), col = "blue", lty = 2, lwd = 2)
```




*# page 26*

```
plot(x, resid(quadratic.lm), ylab = "Residual", xlab = "X", pch = 23, bg = "orange", cex = 2)  
abline(h = 0, lwd = 2, col = "red", lty = 2)
```




```
# page 28  
qqnorm(resid(simple.lm), pch = 23, bg = "orange", cex = 2)  
abline(0, 1)
```

Normal Q-Q Plot




```
qqnorm(resid(quadratic.lm), pch = 23, bg = "orange", cex = 2)  
abline(0, 1)
```

Normal Q-Q Plot




*# page 29*

```
load("/Users/ohsukju/Downloads/data/HIV.VL.Rdata")  
attach(viral.load)  
plot(GSS, VL, pch=23, bg='orange', cex=2)  
viral.lm = lm(VL ~ GSS)  
abline(viral.lm, col='red', lwd=2)
```

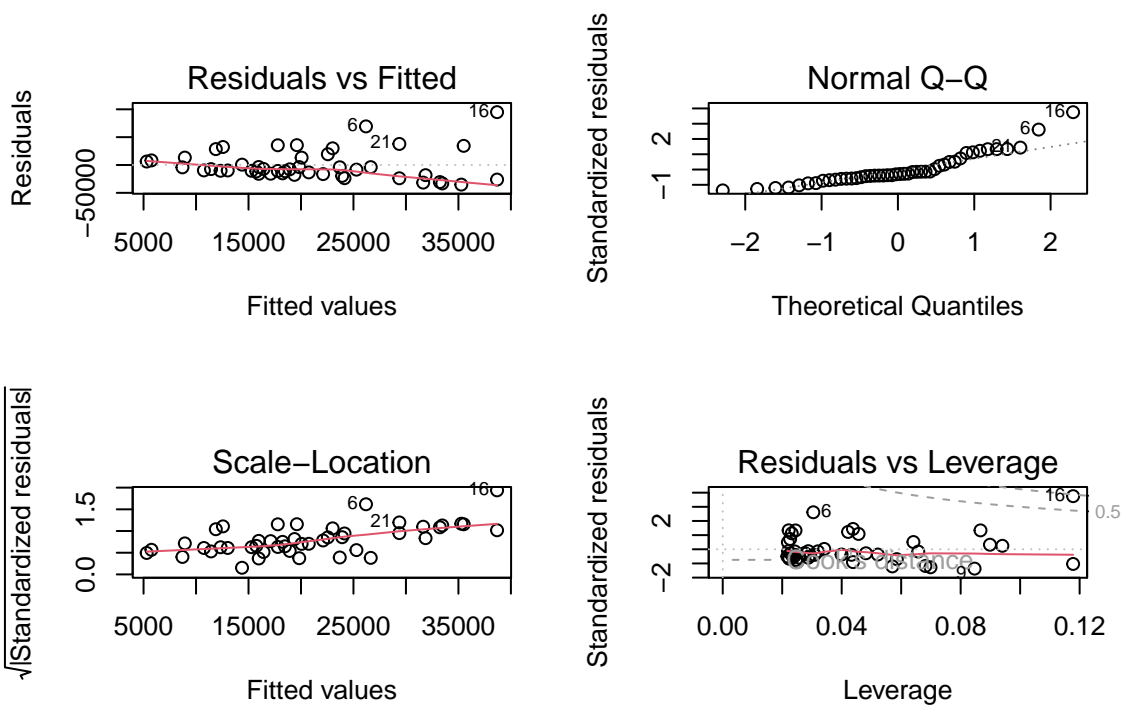


*# page 30*

```
good = (VL < 200000)  
viral.lm.good = lm(VL ~ GSS, subset=good)  
#abline(viral.lm.good, col='blue', lwd=2)  
  
plot(GSS[good], resid(viral.lm.good), pch=23, bg='orange', cex=2, xlab='GSS', ylab='Residual')  
abline(h=0, lwd=2, col='red', lty=2)
```



```
par(mfrow=c(2,2))
plot(viral.lm.good)
```




## 0.2 4. Multiple linear regression

# page 4

```
load("/Users/ohsukju/Downloads/data/prostate.Rdata")

pairs(prostate, pch=23, bg='orange', cex.labels=1.5)
```



# page 7

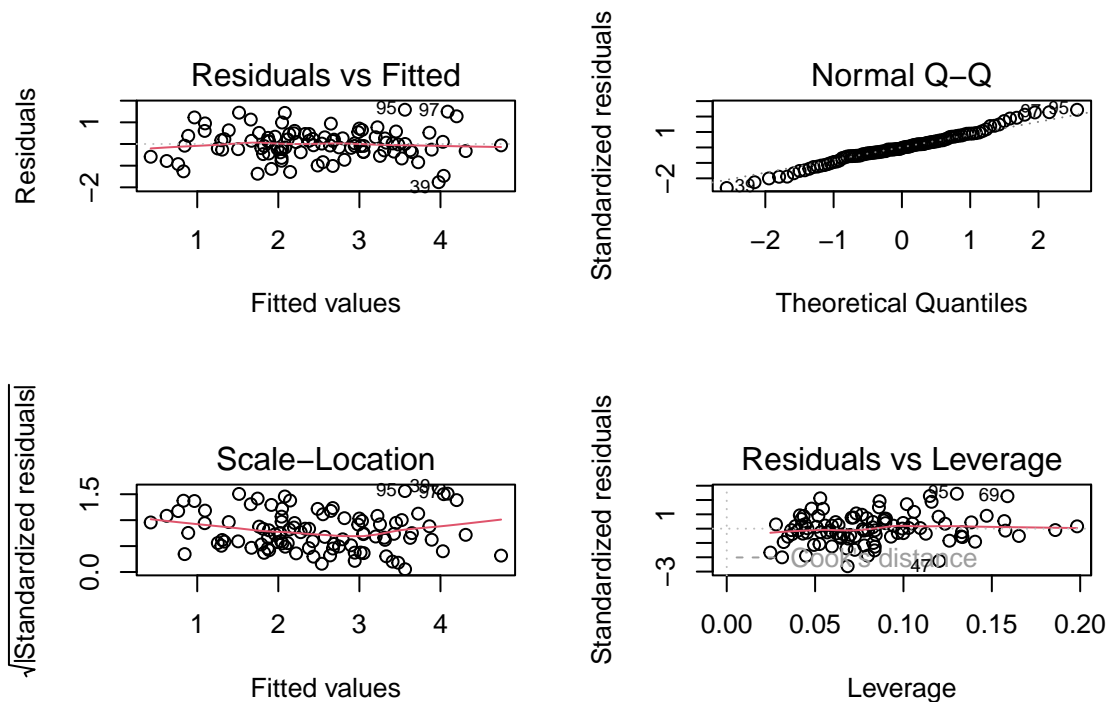
```
prostate.lm = lm(lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45, data=prostate)

summary(prostate.lm)
```

```
##
## Call:
## lm(formula = lpsa ~ lcavol + lweight + age + lbph + svi + lcp +
##      pgg45, data = prostate)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.76395 -0.35764 -0.02143  0.37762  1.58178
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.494155   0.873567   0.566  0.57304
## lcavol       0.569546   0.085847   6.634 2.46e-09 ***
```

```
## lweight      0.614420    0.198449    3.096  0.00262 **
## age         -0.020913    0.010978   -1.905  0.06000 .
## lbph        0.097353    0.057584    1.691  0.09441 .
## svi         0.752397    0.238180    3.159  0.00216 **
## lcp         -0.104959    0.089347   -1.175  0.24323
## pgg45       0.005324    0.003385    1.573  0.11923
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.696 on 89 degrees of freedom
## Multiple R-squared:  0.663, Adjusted R-squared:  0.6365
## F-statistic: 25.01 on 7 and 89 DF,  p-value: < 2.2e-16
```

```
par(mfrow=c(2,2))
plot(prostate.lm)
```



```
# Y_i=b0+b1lcaivol_i+b2lweight_i+b3age_i+b4lbph_i+b5svi_i+b6lcp_i+b7pgg45_i+e_i
```

```
# page 8
```

```
print(prostate.lm$df.resid)
```

```
## [1] 89
```

```
sigma.hat = sqrt(sum(resid(prostate.lm)^2) / prostate.lm$df.resid)
```

```
sigma.hat
```



```
## [1] 0.6959559
```

```
summary(prostate.lm)$sigma
```

```
## [1] 0.6959559
```

```
# page 9
```

```
case1 = prostate[10,]
```

```
case2 = case1
```

```
case2['lcavol'] = case2['lcavol'] + 1
```

```
rbind(case1, case2)
```

```
##          lcavol  lweight age      lbph svi          lcp gleason pgg45      lpsa train
## 10  0.2231436 3.244544  63 -1.386294   0 -1.386294         6      0 1.047319 FALSE
## 101 1.2231436 3.244544  63 -1.386294   0 -1.386294         6      0 1.047319 FALSE
```

```
Yhat = predict(prostate.lm, rbind(case1, case2))
```

```
diff(Yhat)
```

```
##          101
```

```
## 0.569546
```

```
c(Yhat[2]-Yhat[1], coef(prostate.lm)['lcavol'])
```

```
##          101  lcavol
```

```
## 0.569546 0.569546
```

```
# page 10
```

```
partial_resid_lcavol = resid(lm(lcavol ~ lweight + age + lbph + svi + lcp + pgg45, data=prostate))
```

```
partial_resid_lpsa = resid(lm(lpsa ~ lweight + age + lbph + svi + lcp + pgg45, data=prostate))
```

```
summary(lm(partial_resid_lpsa ~ partial_resid_lcavol))
```

```
##
```

```
## Call:
```

```
## lm(formula = partial_resid_lpsa ~ partial_resid_lcavol)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -1.76395 -0.35764 -0.02143  0.37762  1.58178
```

```
##
```

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.956e-17  6.840e-02   0.000      1
## partial_resid_lcavol 5.695e-01  8.309e-02   6.854 7.15e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6736 on 95 degrees of freedom
## Multiple R-squared:  0.3309, Adjusted R-squared:  0.3239
## F-statistic: 46.98 on 1 and 95 DF,  p-value: 7.147e-10
```

```
## simulation
```

```
# generative model
```

```
set.seed(101)
```

```
ns<-100
```

```
#install.packages('mvtnorm')
```

```
library(mvtnorm)
```

```
## Warning: package 'mvtnorm' was built under R version 4.2.3
```

```
rho.x<-0.8
```

```
xs<-rmvnorm(ns, mean=c(0, 0), sigma=matrix(c(1, rho.x, rho.x, 1), 2, 2))
```

```
Xs<-cbind(1, xs)
```

```
head(Xs)
```

```
##      [,1]      [,2]      [,3]
## [1,]    1 -0.04454745  0.3483290
## [2,]    1 -0.50782366 -0.1101151
## [3,]    1  0.80297412  1.1890076
## [4,]    1  0.50304615  0.1758986
## [5,]    1  0.72037041  0.2104183
## [6,]    1  0.11540426 -0.4754957
```

```
tr.bt<-c(5, 2, 4)
```



```
tr.sigma<-1
```

```
ys<-Xs%*%tr.bt+rnorm(ns, 0, tr.sigma)
```

```
par(mfrow=c(1,2))
```

```
plot(Xs[,2], ys)
```

```
plot(Xs[,3], ys)
```

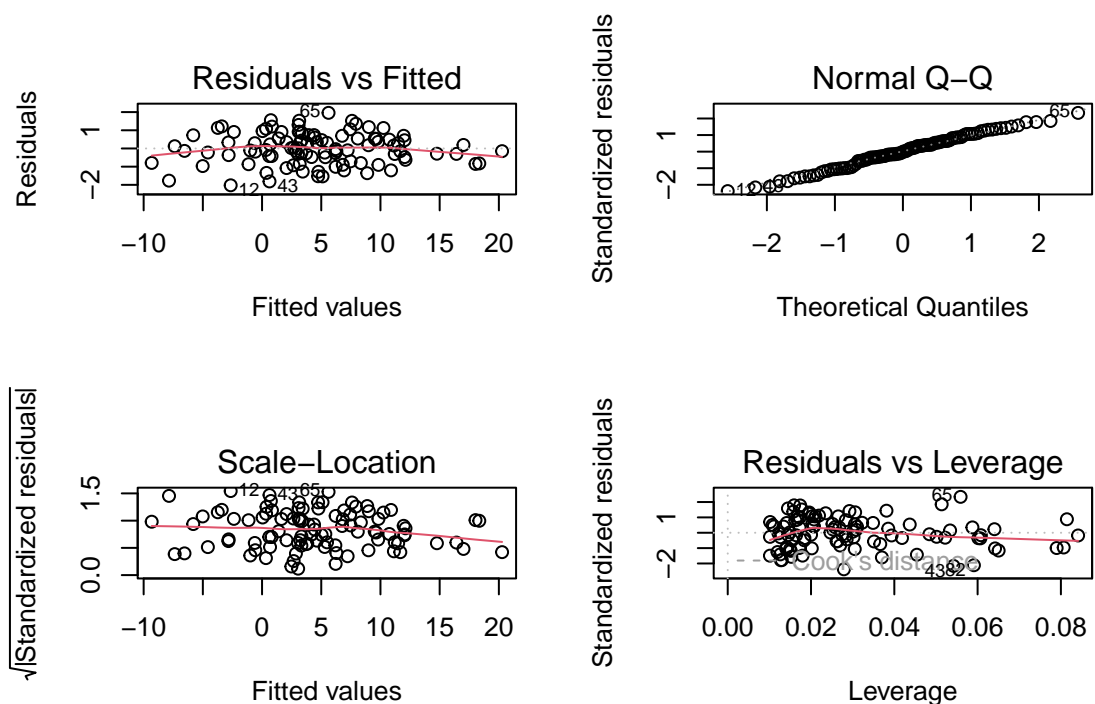


```
datas<-data.frame(y=ys, x1=Xs[,2], x2=Xs[,3])

lm1<-lm(ys~x1+x2, data=datas)

par(mfrow=c(2,2))

plot(lm1)
```



```
summary(lm1)
```

```
##
## Call:
## lm(formula = ys ~ x1 + x2, data = datas)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.02519 -0.51758 -0.01226  0.56301  1.95904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.00099    0.08714   57.39  <2e-16 ***
## x1             2.22736    0.15940   13.97  <2e-16 ***
## x2             3.95810    0.14676   26.97  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8628 on 97 degrees of freedom
## Multiple R-squared:  0.9791, Adjusted R-squared:  0.9787
## F-statistic: 2276 on 2 and 97 DF, p-value: < 2.2e-16
```

```
lm0.1<-lm(ys~x1, data=datas)
```

```
lm0.2<-lm(ys~x2, data=datas)
```

```
# results
```

```
cbind(rbind(summary(lm0.1)$coef[-1,c(1, 4)],  
summary(lm0.2)$coef[-1,c(1, 4)]),  
summary(lm1)$coef[-1,c(1, 4)])
```

```
##      Estimate      Pr(>|t|) Estimate      Pr(>|t|)  
## x1 5.726837 1.360677e-38 2.227357 5.786541e-25  
## x2 5.627316 1.089139e-60 3.958104 7.240269e-47
```

```
# page 11
```

```
Y = prostate$lpsa  
n = length(Y)  
SST = sum((Y - mean(Y))^2)  
MST = SST / (n - 1)  
SSE = sum(resid(prostate.lm)^2)  
MSE = SSE / prostate.lm$df.residual  
SSR = SST - SSE  
MSR = SSR / (n - 1 - prostate.lm$df.residual)  
print(c(MST,MSE,MSR))
```

```
## [1] 1.3324756 0.4843546 12.1157287
```

```
# page 13
```

```
print(summary(prostate.lm))
```

```
##  
## Call:  
## lm(formula = lpsa ~ lcavol + lweight + age + lbph + svi + lcp +  
##      pgg45, data = prostate)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1.76395 -0.35764 -0.02143  0.37762  1.58178   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)  0.494155   0.873567   0.566  0.57304      
## lcavol       0.569546   0.085847   6.634 2.46e-09 ***  
## lweight      0.614420   0.198449   3.096 0.00262 **
```

```
## age          -0.020913   0.010978  -1.905   0.06000 .
## lbph          0.097353   0.057584   1.691   0.09441 .
## svi           0.752397   0.238180   3.159   0.00216 **
## lcp          -0.104959   0.089347  -1.175   0.24323
## pgg45         0.005324   0.003385   1.573   0.11923
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.696 on 89 degrees of freedom
## Multiple R-squared:  0.663, Adjusted R-squared:  0.6365
## F-statistic: 25.01 on 7 and 89 DF,  p-value: < 2.2e-16
```

```
F = MSR / MSE
```

```
F
```

```
## [1] 25.01417
```

```
# page 17
```

```
n = length(Y)
```

```
attach(prostate)
```

```
X = cbind(rep(1,n), lcavol, lweight, age, lbph, svi, lcp, pgg45)
```

```
detach(prostate)
```

```
colnames(X)[1] = '(Intercept)'
```

```
head(X)
```

```
##      (Intercept)      lcavol  lweight age      lbph svi      lcp pgg45
## [1,]           1 -0.5798185  2.769459  50 -1.386294  0 -1.386294    0
## [2,]           1 -0.9942523  3.319626  58 -1.386294  0 -1.386294    0
## [3,]           1 -0.5108256  2.691243  74 -1.386294  0 -1.386294   20
## [4,]           1 -1.2039728  3.282789  58 -1.386294  0 -1.386294    0
## [5,]           1  0.7514161  3.432373  62 -1.386294  0 -1.386294    0
## [6,]           1 -1.0498221  3.228826  50 -1.386294  0 -1.386294    0
```

```
head(model.matrix(prostate.lm))
```

```
##      (Intercept)      lcavol  lweight age      lbph svi      lcp pgg45
## 1           1 -0.5798185  2.769459  50 -1.386294  0 -1.386294    0
## 2           1 -0.9942523  3.319626  58 -1.386294  0 -1.386294    0
## 3           1 -0.5108256  2.691243  74 -1.386294  0 -1.386294   20
## 4           1 -1.2039728  3.282789  58 -1.386294  0 -1.386294    0
## 5           1  0.7514161  3.432373  62 -1.386294  0 -1.386294    0
## 6           1 -1.0498221  3.228826  50 -1.386294  0 -1.386294    0
```

*# page 18*

```
beta = as.numeric(solve(t(X) %*% X) %*% t(X) %*% Y)
```

```
names(beta) = colnames(X)
```

```
print(beta)
```

```
## (Intercept)      lcavol      lweight      age      lbph      svi
## 0.494154754 0.569546032 0.614419817 -0.020913467 0.097352535 0.752397342
##          lcp      pgg45
## -0.104959408 0.005324465
```

```
print(coef(prostate.lm))
```

```
## (Intercept)      lcavol      lweight      age      lbph      svi
## 0.494154754 0.569546032 0.614419817 -0.020913467 0.097352535 0.752397342
##          lcp      pgg45
## -0.104959408 0.005324465
```

*# page 21*

```
confint(prostate.lm, level=0.90)
```

```
##              5 %              95 %
## (Intercept) -0.9578488958  1.946158404
## lcavol      0.4268548240  0.712237239
## lweight     0.2845659251  0.944273708
## age        -0.0391601782 -0.002666755
## lbph       0.0016386253  0.193066445
## svi        0.3565053323  1.148289353
## lcp        -0.2534678904  0.043549074
## pgg45      -0.0003011464  0.010950077
```

*# page 22*

```
summary(prostate.lm)$coef
```

```
##              Estimate Std. Error  t value    Pr(>|t|)
## (Intercept)  0.494154754 0.873566764  0.5656749 5.730382e-01
## lcavol      0.569546032 0.085847096  6.6344240 2.461450e-09
## lweight     0.614419817 0.198449499  3.0961016 2.622121e-03
## age        -0.020913467 0.010977742 -1.9050792 5.999883e-02
## lbph       0.097352535 0.057584215  1.6906115 9.441057e-02
## svi        0.752397342 0.238179913  3.1589454 2.163265e-03
## lcp        -0.104959408 0.089346934 -1.1747399 2.432323e-01
## pgg45      0.005324465 0.003384528  1.5731780 1.192260e-01
```

*# H0: bt1=0 vs HA: bt1 != 0*

```
T1 = 0.56954 / 0.08584
```

```
P1 = 2 * (1 - pt(abs(T1), 89))
```

```
print(c(T1,P1))
```

```
## [1] 6.634902e+00 2.456119e-09
```

```
# page 24
```

```
Y.hat = X %*% beta
```

```
sigma.hat = sqrt(sum((Y - Y.hat)^2) / (n - ncol(X)))
```

```
cov.beta = sigma.hat^2 * solve(t(X) %*% X)
```

```
cov.beta
```

```
##      (Intercept)      lcavol      lweight      age
## (Intercept)  0.763118892  9.968185e-03 -1.127145e-01 -5.757711e-03
## lcavol      0.009968185  7.369724e-03 -3.267921e-03 -1.303787e-04
## lweight     -0.112714488 -3.267921e-03  3.938220e-02 -4.088331e-04
## age        -0.005757711 -1.303787e-04 -4.088331e-04  1.205108e-04
## lbph       0.024821341  3.981644e-04 -4.540916e-03 -1.471652e-04
## svi        0.010342177 -2.643164e-03 -4.105065e-03 -7.340636e-05
## lcp        0.002077485 -3.502258e-03  6.198759e-05  1.429198e-04
## pgg45      0.000033477  7.077893e-06  6.876251e-05 -9.141775e-06
##      lbph      svi      lcp      pgg45
## (Intercept)  2.482134e-02  1.034218e-02  2.077485e-03  3.347700e-05
## lcavol      3.981644e-04 -2.643164e-03 -3.502258e-03  7.077893e-06
## lweight     -4.540916e-03 -4.105065e-03  6.198759e-05  6.876251e-05
## age        -1.471652e-04 -7.340636e-05  1.429198e-04 -9.141775e-06
## lbph       3.315942e-03  2.136331e-03 -1.416010e-04 -1.238514e-05
## svi        2.136331e-03  5.672967e-02 -8.863127e-03 -5.227255e-05
## lcp        -1.416010e-04 -8.863127e-03  7.982875e-03 -1.368737e-04
## pgg45      -1.238514e-05 -5.227255e-05 -1.368737e-04  1.145503e-05
```

```
sqrt(diag(cov.beta))
```

```
## (Intercept)      lcavol      lweight      age      lbph      svi
## 0.873566764 0.085847096 0.198449499 0.010977742 0.057584215 0.238179913
##      lcp      pgg45
## 0.089346934 0.003384528
```

```
vcov(prostate.lm)
```

```
##      (Intercept)      lcavol      lweight      age
## (Intercept)  0.763118892  9.968185e-03 -1.127145e-01 -5.757711e-03
## lcavol      0.009968185  7.369724e-03 -3.267921e-03 -1.303787e-04
## lweight     -0.112714488 -3.267921e-03  3.938220e-02 -4.088331e-04
## age        -0.005757711 -1.303787e-04 -4.088331e-04  1.205108e-04
```



```
## lbph      0.024821341  3.981644e-04 -4.540916e-03 -1.471652e-04
## svi       0.010342177 -2.643164e-03 -4.105065e-03 -7.340636e-05
## lcp       0.002077485 -3.502258e-03  6.198759e-05  1.429198e-04
## pgg45     0.000033477  7.077893e-06  6.876251e-05 -9.141775e-06
##          lbph      svi      lcp      pgg45
## (Intercept) 2.482134e-02  1.034218e-02  2.077485e-03  3.347700e-05
## lcavol      3.981644e-04 -2.643164e-03 -3.502258e-03  7.077893e-06
## lweight    -4.540916e-03 -4.105065e-03  6.198759e-05  6.876251e-05
## age        -1.471652e-04 -7.340636e-05  1.429198e-04 -9.141775e-06
## lbph       3.315942e-03  2.136331e-03 -1.416010e-04 -1.238514e-05
## svi        2.136331e-03  5.672967e-02 -8.863127e-03 -5.227255e-05
## lcp       -1.416010e-04 -8.863127e-03  7.982875e-03 -1.368737e-04
## pgg45     -1.238514e-05 -5.227255e-05 -1.368737e-04  1.145503e-05
```

*# page 25*

```
CI.lm = function(cur.lm, a, level=0.95, extra=0) {

  # the center of the confidence interval

  center = sum(a*cur.lm$coef)

  # the estimate of sigma^2

  sigma.sq = sum(resid(cur.lm)^2) / cur.lm$df.resid

  # the standard error of sum(a*cur.lm$coef)

  se = sqrt(extra * sigma.sq + sum((a %*% vcov(cur.lm)) * a))

  # the degrees of freedom for the t-statistic

  df = cur.lm$df

  # the quantile used in the confidence interval

  q = qt((1 - level)/2, df, lower.tail=FALSE)

  # upper, lower limits

  upr = center + se * q

  lwr = center - se * q

  return(data.frame(center, lwr, upr))

}
```

```
# 95% CI for E(y|xs) with xs=(lcavol=1.3, lweight=3.6, age=64, lbph=0.1, svi=0.2, lcp=-0.2, pgg45=25)

print(CI.lm(prostate.lm, c(1, 1.3, 3.6, 64, 0.1, 0.2, -0.2, 25)))
```

```
##      center      lwr      upr
## 1 2.422332 2.281218 2.563447
```

```
predict(prostate.lm, list(lcavol=1.3, lweight=3.6, age=64, lbph=0.1, svi=0.2, lcp=-0.2, pgg45=25), interval=
```

```
##          fit          lwr          upr
## 1 2.422332 2.281218 2.563447
```

```
# page 26
```

```
CI.lm(prostate.lm, c(0,0,0,2,0,0.5,0,0))
```

```
##          center          lwr          upr
## 1 0.3343717 0.09496226 0.5737812
```

```
case1 = prostate[10,]
case2 = case1
case2['age'] = case2['age'] + 2
case2['svi'] = case2['svi'] + 0.5
rbind(case1, case2)
```

```
##          lcavol lweight age          lbph svi          lcp gleason pgg45          lpsa train
## 10 0.2231436 3.244544 63 -1.386294 0.0 -1.386294          6          0 1.047319 FALSE
## 101 0.2231436 3.244544 65 -1.386294 0.5 -1.386294          6          0 1.047319 FALSE
```

```
Yhat = predict(prostate.lm, rbind(case1, case2))
diff(Yhat)
```

```
##          101
## 0.3343717
```

```
# page 29
```

```
f.test.lm = function(R.lm, F.lm) {
  SSE.R = sum(resid(R.lm)^2)
  SSE.F = sum(resid(F.lm)^2)
  df.num = R.lm$df - F.lm$df
  df.den = F.lm$df
  F = ((SSE.R - SSE.F) / df.num) / (SSE.F / df.den)
  p.value = 1 - pf(F, df.num, df.den)
  return(data.frame(F, df.num, df.den, p.value))
}
```

```

}

reduced.lm = lm(lpsa ~ lcavol + lweight + age + lbph + svi, data=prostate)

print(f.test.lm(reduced.lm, prostate.lm))

##           F df.num df.den   p.value
## 1 1.372057      2      89 0.2588958

anova(reduced.lm, prostate.lm)

## Analysis of Variance Table
##
## Model 1: lpsa ~ lcavol + lweight + age + lbph + svi
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      91 44.437
## 2      89 43.108  2    1.3291 1.3721 0.2589

anova(prostate.lm, reduced.lm)

## Analysis of Variance Table
##
## Model 1: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      89 43.108
## 2      91 44.437 -2   -1.3291 1.3721 0.2589

# page 32

prostate$Z = prostate$lcavol + prostate$svi

equal.lm = lm(lpsa ~ Z + lweight + age + lbph + lcp + pgg45, data=prostate)

f.test.lm(equal.lm, prostate.lm)

##           F df.num df.den   p.value
## 1 0.4818657      1      89 0.4893864

anova(equal.lm, prostate.lm)

## Analysis of Variance Table
##
## Model 1: lpsa ~ Z + lweight + age + lbph + lcp + pgg45
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      90 43.341
## 2      89 43.108  1    0.23339 0.4819 0.4894

```

```
anova(prostate.lm)
```

```
## Analysis of Variance Table
##
## Response: lpsa
##           Df Sum Sq Mean Sq  F value    Pr(>F)
## lcavol     1  69.003   69.003  142.4636 < 2.2e-16 ***
## lweight     1   7.173    7.173   14.8086 0.0002235 ***
## age         1   0.646    0.646    1.3330 0.2513641
## lbph        1   0.809    0.809    1.6702 0.1995770
## svi         1   5.851    5.851   12.0798 0.0007900 ***
## lcp         1   0.130    0.130    0.2692 0.6051404
## pgg45       1   1.199    1.199    2.4749 0.1192260
## Residuals  89  43.108    0.484
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# page 33
```

```
prostate$Z2 = prostate$lcavol - prostate$svi
prostate$Y2 = prostate$lpsa - prostate$svi
constrained.lm = lm(lpsa ~ Z2 + lweight + age + lbph + lcp + pgg45, data=prostate, offset=svi)
constrained.lm2 = lm(Y2 ~ Z2 + lweight + age + lbph + lcp + pgg45, data=prostate)
coef(constrained.lm)
```

```
## (Intercept)          Z2      lweight          age          lbph          lcp
## 0.382975608 0.543672792 0.654779620 -0.019797945 0.083478678 -0.037271152
##           pgg45
## 0.005571861
```

```
coef(constrained.lm2)
```

```
## (Intercept)          Z2      lweight          age          lbph          lcp
## 0.382975608 0.543672792 0.654779620 -0.019797945 0.083478678 -0.037271152
##           pgg45
## 0.005571861
```

```
anova(constrained.lm, prostate.lm)
```

```
## Analysis of Variance Table
##
## Model 1: lpsa ~ Z2 + lweight + age + lbph + lcp + pgg45
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      90 43.961
## 2      89 43.108  1  0.85359 1.7623 0.1877
```

```
anova(constrained.lm2, prostate.lm)
```

```
## Warning in anova.lm(object, ...): models with response '"lpsa"' removed
## because response differs from model 1
```

```
## Analysis of Variance Table
##
```

```
## Response: Y2
```

```
##           Df Sum Sq Mean Sq F value    Pr(>F)
## Z2          1 37.663   37.663  77.1051 9.831e-14 ***
## lweight     1  7.619    7.619  15.5975 0.0001554 ***
## age         1  0.569    0.569   1.1659 0.2831319
## lbph        1  1.129    1.129   2.3109 0.1319718
## lcp         1  0.161    0.161   0.3296 0.5673017
## pgg45       1  1.317    1.317   2.6956 0.1041138
## Residuals  90 43.961    0.488
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
C.mat = matrix(0, 2, 8)
```

```
C.mat[1,7] = 1
```

```
C.mat[2,8] = 1
```

```
print(C.mat)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
## [1,]    0    0    0    0    0    0    1    0
## [2,]    0    0    0    0    0    0    0    1
```

```
#install.packages('car')
```

```
library(car)
```

```
## Loading required package: carData
```

```
linearHypothesis(prostate.lm, C.mat, test='F')
```

```
## Linear hypothesis test
```

```
##
```

```
## Hypothesis:
```

```
## lcp = 0
```

```
## pgg45 = 0
```

```
##
```

```
## Model 1: restricted model
```

```
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45
```

```
##
```

```
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
```

```
## 1      91 44.437
```

```
## 2      89 43.108  2    1.3291 1.3721 0.2589
```

```
anova(reduced.lm, prostate.lm)
```

```
## Analysis of Variance Table
##
## Model 1: lpsa ~ lcavol + lweight + age + lbph + svi
## Model 2: lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      91 44.437
## 2      89 43.108  2    1.3291 1.3721 0.2589
```

```
head(prostate)
```

```
##      lcavol lweight age      lbph svi      lcp gleason pgg45      lpsa
## 1 -0.5798185 2.769459 50 -1.386294  0 -1.386294      6      0 -0.4307829
## 2 -0.9942523 3.319626 58 -1.386294  0 -1.386294      6      0 -0.1625189
## 3 -0.5108256 2.691243 74 -1.386294  0 -1.386294      7     20 -0.1625189
## 4 -1.2039728 3.282789 58 -1.386294  0 -1.386294      6      0 -0.1625189
## 5  0.7514161 3.432373 62 -1.386294  0 -1.386294      6      0  0.3715636
## 6 -1.0498221 3.228826 50 -1.386294  0 -1.386294      6      0  0.7654678
##   train      Z      Z2      Y2
## 1  TRUE -0.5798185 -0.5798185 -0.4307829
## 2  TRUE -0.9942523 -0.9942523 -0.1625189
## 3  TRUE -0.5108256 -0.5108256 -0.1625189
## 4  TRUE -1.2039728 -1.2039728 -0.1625189
## 5  TRUE  0.7514161  0.7514161  0.3715636
## 6  TRUE -1.0498221 -1.0498221  0.7654678
```

```
prostate2<-prostate[,1:9]
```

```
head(prostate2)
```

```
##      lcavol lweight age      lbph svi      lcp gleason pgg45      lpsa
## 1 -0.5798185 2.769459 50 -1.386294  0 -1.386294      6      0 -0.4307829
## 2 -0.9942523 3.319626 58 -1.386294  0 -1.386294      6      0 -0.1625189
## 3 -0.5108256 2.691243 74 -1.386294  0 -1.386294      7     20 -0.1625189
## 4 -1.2039728 3.282789 58 -1.386294  0 -1.386294      6      0 -0.1625189
## 5  0.7514161 3.432373 62 -1.386294  0 -1.386294      6      0  0.3715636
## 6 -1.0498221 3.228826 50 -1.386294  0 -1.386294      6      0  0.7654678
```

```
prostate2<-scale(prostate2)
```

```
summary(prostate2)
```

```
##      lcavol      lweight      age      lbph
## Min.   :-2.28833  Min.   :-2.92718  Min.   :-3.0713  Min.   :-1.0247
## 1st Qu.: -0.71031  1st Qu.: -0.59070  1st Qu.: -0.5193  1st Qu.: -1.0247
## Median :  0.08222  Median : -0.01385  Median :  0.1523  Median :  0.1377
## Mean    :  0.00000  Mean    :  0.00000  Mean    :  0.0000  Mean    :  0.0000
## 3rd Qu.:  0.65927  3rd Qu.:  0.57761  3rd Qu.:  0.5553  3rd Qu.:  1.0048
## Max.    :  2.09651  Max.    :  2.68770  Max.    :  2.0327  Max.    :  1.5343
##      svi      lcp      gleason      pgg45
```

```
## Min.      :-0.5229   Min.      :-0.8632   Min.      :-1.0422   Min.      :-0.8645
## 1st Qu.: -0.5229   1st Qu.: -0.8632   1st Qu.: -1.0422   1st Qu.: -0.8645
## Median : -0.5229   Median : -0.4428   Median :  0.3426   Median : -0.3326
## Mean    :  0.0000   Mean    :  0.0000   Mean    :  0.0000   Mean    :  0.0000
## 3rd Qu.: -0.5229   3rd Qu.:  0.9712   3rd Qu.:  0.3426   3rd Qu.:  0.5538
## Max.    :  1.8925   Max.    :  2.2053   Max.    :  3.1122   Max.    :  2.6811
##      lpsa
## Min.      :-2.5202
## 1st Qu.: -0.6469
## Median :  0.0980
## Mean    :  0.0000
## 3rd Qu.:  0.5007
## Max.    :  2.6895
```

```
apply(prostate2, 2, sd)
```

```
## lcavol lweight    age    lbph    svi    lcp gleason    pgg45    lpsa
##      1      1      1      1      1      1      1      1      1
```

```
prostate.lm2 = lm(lpsa ~ lcavol + lweight + age + lbph + svi + lcp + pgg45, data=as.data.frame(prostate2))
```

```
summary(prostate.lm)$coef
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept)  0.494154754 0.873566764  0.5656749 5.730382e-01
## lcavol       0.569546032 0.085847096  6.6344240 2.461450e-09
## lweight      0.614419817 0.198449499  3.0961016 2.622121e-03
## age          -0.020913467 0.010977742 -1.9050792 5.999883e-02
## lbph         0.097352535 0.057584215  1.6906115 9.441057e-02
## svi          0.752397342 0.238179913  3.1589454 2.163265e-03
## lcp          -0.104959408 0.089346934 -1.1747399 2.432323e-01
## pgg45         0.005324465 0.003384528  1.5731780 1.192260e-01
```

```
summary(prostate.lm2)$coef
```

```
##              Estimate Std. Error    t value    Pr(>|t|)
## (Intercept) -8.348482e-16 0.06121618 -1.363771e-14 1.000000e+00
## lcavol       5.815336e-01 0.08765397  6.634424e+00 2.461450e-09
## lweight      2.280323e-01 0.07365144  3.096102e+00 2.622121e-03
## age          -1.348863e-01 0.07080353 -1.905079e+00 5.999883e-02
## lbph         1.223565e-01 0.07237413  1.690611e+00 9.441057e-02
## svi          2.698439e-01 0.08542214  3.158945e+00 2.163265e-03
## lcp          -1.271383e-01 0.10822678 -1.174740e+00 2.432323e-01
## pgg45         1.300941e-01 0.08269510  1.573178e+00 1.192260e-01
```

*#Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code to the console*